

Classification and Regression Trees (CART)

Regression Trees: Ozone data

This example is adapted from the book “Extending the Linear Model with R”, by Julian J. Faraway. Here is a quote from that book describing the data:

We apply the regression tree methodology to study the relationship between atmospheric ozone concentration and meteorology in the Los Angeles Basin in 1976. A number of cases with missing variables hve been removed for simplicity [but Evan notes that trees are among the regression and classification methods that are best able to handle missing data]. We wish to predict the ozone level from the other predictors.

The variables in the data set are as follows:

- `o3`: Ozone concentration (ppm) at Sandbug Air Force Base
- `vh`: Vandenburg 500 millibar height (inches)
- `wind`: wind speed (miles per hour)
- `humidity`: humidity (percent)
- `temp`: temperature (degrees C)
- `ibh`: inversion base height (feet)
- `dpg`: Daggett pressure gradient (mmhg)
- `ibt`: inversion base temperature (degrees F)
- `vis`: visibility (miles)
- `doy`: day of the year

```
library(faraway)

##
## Attaching package: 'faraway'

## The following object is masked from 'package:GGally':
##
##      happy


head(ozone)

##      O3      vh wind humidity temp  ibh dpg ibt vis doy
## 1   3 5710     4      28   40 2693 -25  87 250  33
## 2   5 5700     3      37   45  590 -24 128 100  34
## 3   5 5760     3      51   54 1450  25 139  60  35
## 4   6 5720     4      69   35 1568  15 121  60  36
## 5   4 5790     6      19   45 2631 -33 123 100  37
## 6   4 5790     3      25   55  554 -28 182 250  38


dim(ozone)

## [1] 330  10

ggpairs(ozone %>% select(vh:doy, O3))
```



```
ggplot(data = ozone, mapping = aes(x = temp, y = ibh, color = O3)) +
  geom_point() +
  scale_color_gradient2(midpoint = median(ozone$O3)) +
  xlim(c(0, 100))
```



Regression tree with 2 variables:

```
# Load rpart package (stands for "Recursive Partitioning")
```

```
library(rpart)
```

```
##
```

```
## Attaching package: 'rpart'
```

```
## The following object is masked from 'package:faraway':
```

```
##
```

```
## solder
```

```
# fit regression tree model
```

```
roz <- rpart(O3 ~ temp + ibh, data = ozone)
```

```
print(roz)
```

```
## n= 330
```

```
##
```

```
## node), split, n, deviance, yval
```

```
## * denotes terminal node
```

```
##
```

```
## 1) root 330 21115.4100 11.775760
```

```
## 2) temp< 67.5 214 4114.3040 7.425234
```

```
## 4) ibh>=3573.5 108 689.6296 5.148148 *
```

```
## 5) ibh< 3573.5 106 2294.1230 9.745283
```

```
## 10) temp< 58.5 53 1008.5280 7.905660 *
```

```
## 11) temp>=58.5 53 926.8679 11.584910 *
```

```
## 3) temp>=67.5 116 5478.4400 19.801720
```

```
## 6) temp< 79.5 72 2626.0000 17.166670
```

```
## 12) ibh>=2785 15 348.9333 12.066670 *
```

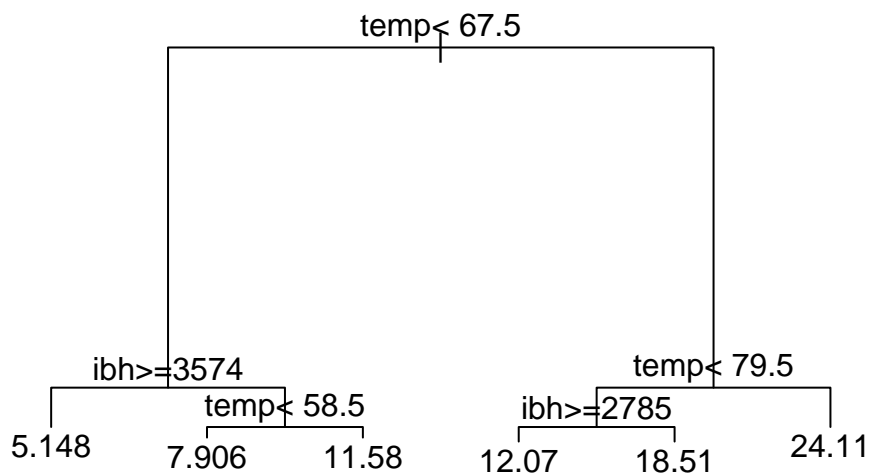
```
## 13) ibh< 2785 57 1784.2460 18.508770 *
```

```
## 7) temp>=79.5 44 1534.4320 24.113640 *
```

```
# print first picture of resulting tree
```

```
plot(roz, margin = 0.1)
```


```
text(roz)
```



```
# print second picture of resulting tree
```

```
plot(roz, margin = 0.1, uniform = TRUE)
```

```
text(roz)
```



What's the predicted ozone level for a day with a temperature of 75 degrees and an inversion base height of 2000 feet?

```
test_data <- data.frame(
  temp = 75, ibh = 2000
)

predict(roz, newdata = test_data)
```

```
##      1
## 18.50877
```

Notation/Mathematical Description:


$$\hat{f}(x_i) = \sum_{m=1}^{|T|} I_{R_m}(x_i) \hat{y}_m$$

- $|T|$ is the number of *terminal nodes* in the tree.
- R_m is the set of values for explanatory variables that fall into the m th terminal node. (R for region or rectangle)
- $I_{R_m}(x_i) = \begin{cases} 1 & \text{if } x_i \text{ is in region } R_m \\ 0 & \text{otherwise} \end{cases}$
- \hat{y}_m is the predicted value in terminal node number m .

More covariates:

```
roz_2 <- rpart(O3 ~ ., data = ozone)

# print picture of resulting tree
plot(roz_2, margin = 0.1, uniform = TRUE)
text(roz_2)
```



Classification Trees: Heart Disease data

We have data on 303 patients who presented with chest pain. The response variable is AHD, which is “Yes” if an angiographic test indicates presence of heart disease, and “No” otherwise. There are 13 predictor variables which are a mix of quantitative and categorical variables.

```
library(readr)
heart <- read_csv("http://www.evanlray.com/data/islr/Heart.csv") %>%
  select(-X1) %>% # drop leading column of row numbers
  mutate_at(c("Sex", "ChestPain", "Fbs", "RestECG", "ExAng", "Slope", "Thal", "AHD"), factor)
```


```
## Warning: Missing column names filled in: 'X1' [1]
```

```
head(heart)
```

```
## # A tibble: 6 x 14
##   Age Sex ChestPain RestBP Chol Fbs RestECG MaxHR ExAng Oldpeak
##   <int> <fct> <fct>      <int> <int> <fct> <fct>    <int> <fct>    <dbl>
## 1   63 1    typical      145   233 1     2      150 0      2.3
## 2   67 1    asymptom~    160   286 0     2      108 1      1.5
## 3   67 1    asymptom~    120   229 0     2      129 1      2.6
## 4   37 1    nonangin~    130   250 0     0      187 0      3.5
## 5   41 0    nontypic~    130   204 0     2      172 0      1.4
## 6   56 1    nontypic~    120   236 0     0      178 0      0.8
## # ... with 4 more variables: Slope <fct>, Ca <int>, Thal <fct>, AHD <fct>
```

```
rhd <- rpart(AHD ~ ., data = as.data.frame(heart))
```

```
# print second picture of resulting tree
plot(rhd, margin = 0.1, uniform = TRUE)
text(rhd)
```



```
levels(heart$Thal)
```

```
## [1] "fixed"      "normal"     "reversible"
```

```
levels(heart$ChestPain)
```

```
## [1] "asymptomatic" "nonanginal"   "nontypical"   "typical"
```

What's the predicted class for someone whose Thallium stress test results are normal, whose chest pain symptoms are typical, and has a “Ca” (not sure what that stands for) of 2?

```
head(person_to_classify)
```

```
##   Age Sex ChestPain RestBP Chol Fbs RestECG MaxHR ExAng Oldpeak Slope Ca
## 1  NA <NA>  typical    NA   NA <NA>    <NA>    NA <NA>    NA <NA>  2
##   Thal AHD
## 1 normal <NA>
```

```
predict(rhd, newdata = person_to_classify, type = "class")
```

```
## 1
## Yes
## Levels: No Yes
```

Overview of Estimation

Reminder of Notation

Recall our mathematical formulation:

$$\hat{f}(x_i) = \sum_{m=1}^{|T|} I_{R_m}(x_i) \hat{y}_m$$

- $|T|$ is the number of *terminal nodes* in the tree.
- R_m is the set of values for explanatory variables that fall into the m th terminal node. (R for region or rectangle)
- $I_{R_m}(x_i) = \begin{cases} 1 & \text{if } x_i \text{ is in region } R_m \\ 0 & \text{otherwise} \end{cases}$
- \hat{y}_m is the predicted value in terminal node number m .

Parameters to Estimate

- **Split Points:** for each branch of the tree, what covariate is used and at which value does the split occur? This determines the R_m .
- **Regression Constants:** In each terminal node, what is the predicted value \hat{y}_m ?

Optimization target for regression

$$RSS = \sum_{i=1}^n (\hat{y}_i - y_i)^2 = \sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_m)^2$$

Optimization target for classification

We could use classification error rate, but in practice other options are used more frequently.

For node m , let p_{jm} be the proportion of the observations in node m that have class j .

One common choice is the **Gini index**: $1 - \sum_J p_{jm}^2$

Top-down Estimation Algorithm (discussed for regression; classification similar)

1. Start with a tree with only one region/terminal node (the same prediction is made for all observations).
 - Predicted value is mean of all observations
 - Calculate the RSS based on that prediction
2. Repeat the following until a stopping criterion is met:
 - a. For every terminal node in the current tree, consider all possible split points for each covariate X_1, \dots, X_p .
 - For each possible split, predicted values will be the mean of all observations in the newly created leaves
 - Calculate RSS based on that split
 - b. Select the split that achieved the lowest RSS

Commonly used stopping criteria:

- All regions contain 5 or fewer observations (for example)
- A maximum number of terminal nodes has been reached
- No improvement in RSS larger than ε can be realized

Regularizing number of terminal nodes

A smaller number of terminal nodes is less likely to overfit. Choose tree to minimize:

$$RSS + \lambda|T|$$

Use cross-validation to select λ .